

# INSIDE THE JAVA VIRTUAL MACHINE

## *Memory Management and Troubleshooting*

Filip Hanik  
Covalent Technologies  
August 29, 2007

# Who am I?

- ✔ [fhanik@apache.org](mailto:fhanik@apache.org)
- ✔ Tomcat Committer / ASF member
- ✔ Co-designed the Comet implementation
- ✔ Implemented NIO connector in 6
- ✔ Responsible for session replication and clustering
- ✔ Been involved with ASF since 2001
- ✔ Member, Covalent Technical Team

# What are we Talking About?

- Internals of Java Memory
- Spoken from a Java developer's standpoint
- For other Java developers and system administrators

# Agenda

- ✔ Understanding the Java Memory Layout
- ✔ Out Of Memory Errors
  - ✔ Causes
  - ✔ Solution
- ✔ Garbage Collection Basics
- ✔ Java Tuning Options – Time Constraint
- ✔ Questions and Answers

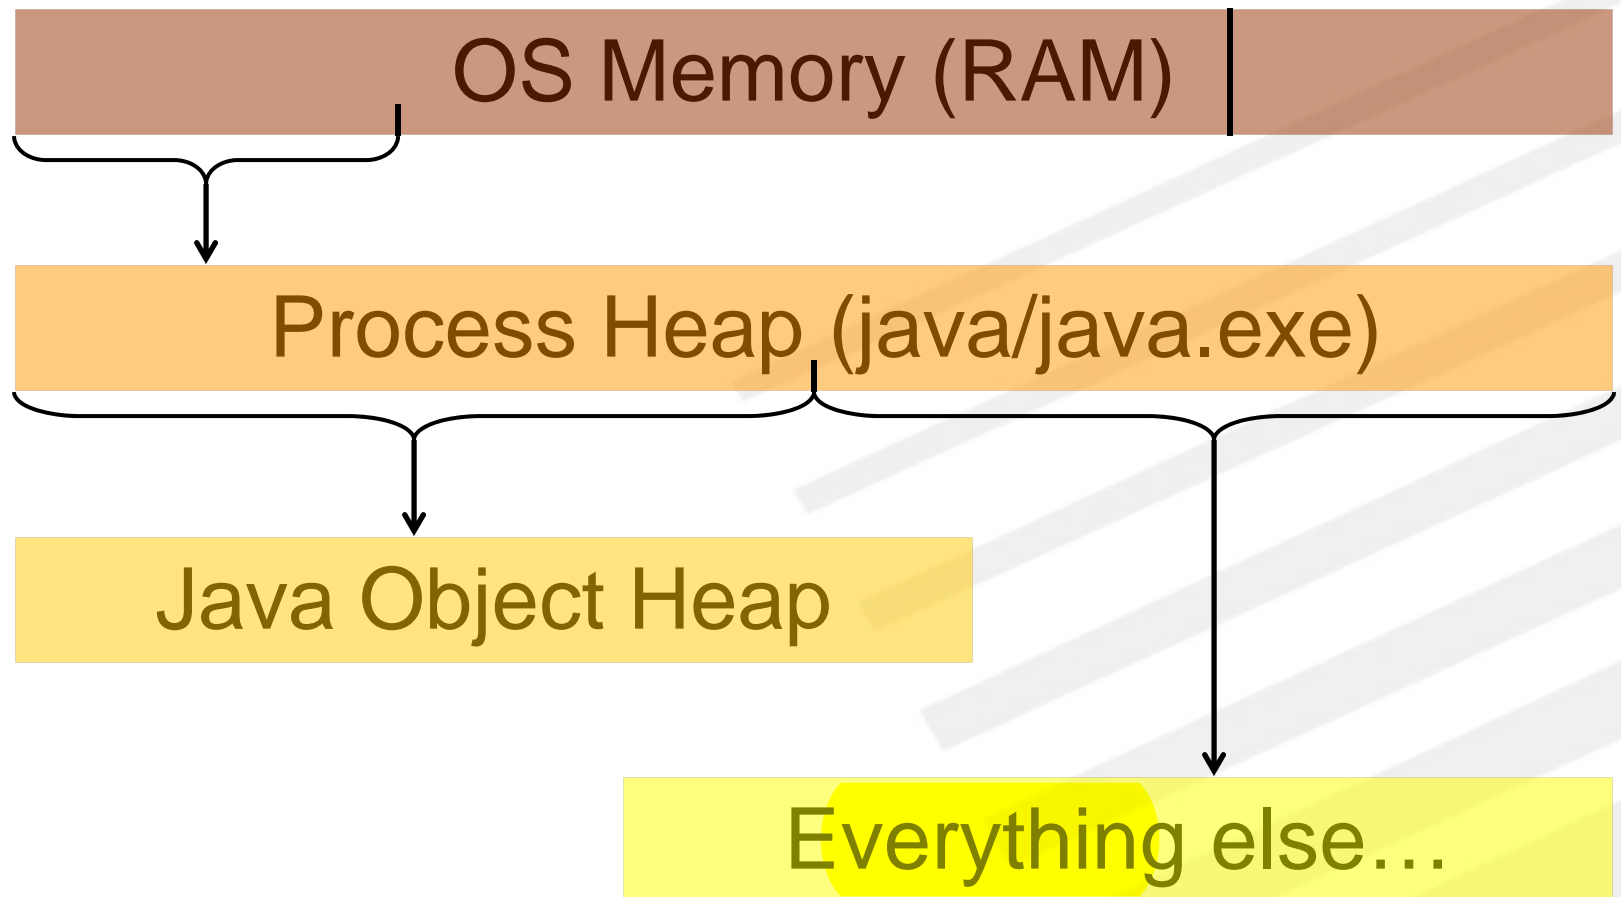
# Storing Data in Memory

- ❖ Java runs as a single process
  - ❖ Does not share memory with other processes
- ❖ Each process allocates memory
  - ❖ We call this **process heap**
- ❖ Ways to allocate memory in a process
  - ❖ C (malloc and free)
  - ❖ C++ (new and delete)
  - ❖ Java (new and dereference -> Garbage Collection)

# Storing Data in Memory

- ✔ JVM manages the process heap
  - ✔ In most cases
  - ✔ JNI managed memory would be an exception, and there are others
- ✔ No shared memory between processes
  - ✔ At least not available through the Java API
- ✔ JVM creates a **Java Heap**
  - ✔ Part of the process heap
  - ✔ Configured through `-Xmx` and `-Xms` settings

# The JVM Process Heap



# JVM Process Heap

- ❖ Maximum size is limited
  - ❖ 32 bit size, roughly 2GB
  - ❖ 64 bit, much much larger 😊
- ❖ If 2GB is the max for the process
  - ❖ -Xmx1800m –Xms1800m – not very good
  - ❖ Leaves no room for anything else



# Java Object Heap

- ✦ Also referred to as Java Heap
  - ✦ Often confused with JVM process heap
- ✦ Stores Java Objects
  - ✦ instances of classes
  - ✦ and the data the objects contain
    - ✦ Primitives
    - ✦ References

# Benefits of the Java Heap

- ✔ Pre-allocate large blocks of memory
- ✔ Allocation of small amounts of memory is very fast
- ✔ No need to fish for a free memory segment in RAM
- ✔ No fragmentation
- ✔ Continuous memory blocks can make a big difference
- ✔ NullPointerException vs. General Access Fault
  - ✔ NPE runtime error
  - ✔ GAF crash the process

# Gotcha #1

- ❖ -Xmx, -Xms and -Xmn
  - ❖ Only controls the Java Object Heap
  - ❖ Often misunderstood to control the process heap
- ❖ Confusion leads to incorrect tuning
  - ❖ And in some cases, the situation worsens

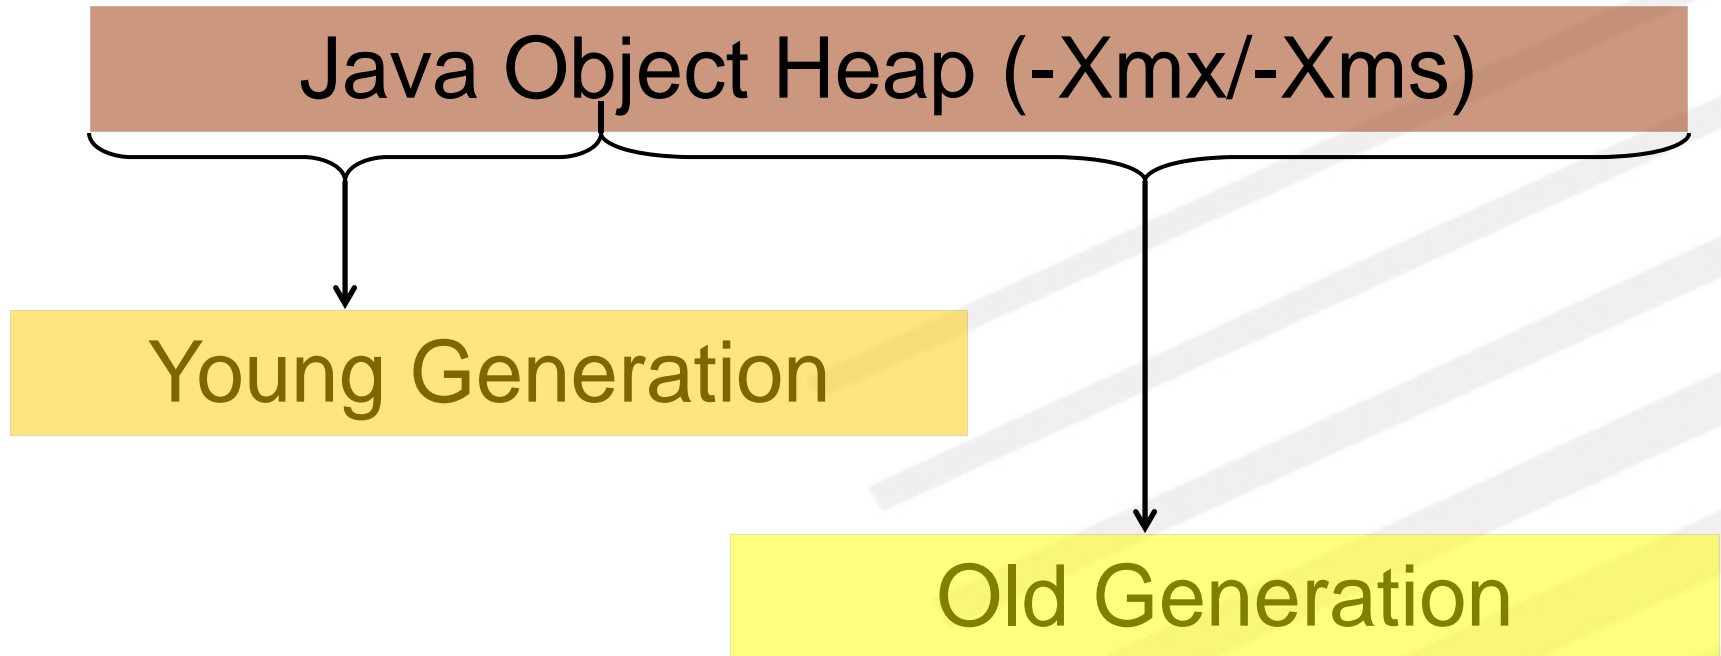
# Java Object Heap

- ✦ So how is the Java Heap allocated?
- ✦ -XX:MinHeapFreeRatio=
  - ✦ Default is 40 (40%)
  - ✦ When the JVM allocates memory, it allocates enough to get 40% free
  - ✦ Huge chunks, very large default
  - ✦ Not important when `-Xms == -Xmx`
- ✦ -XX:MaxHeapFreeRatio=
  - ✦ Default 70%
  - ✦ To avoid over allocation
  - ✦ To give back memory not used
- ✦ As you can see, to provide performance and avoid fragmentation, excessively large blocks are allocated each time

# Java Object Heap

- ❖ Object allocation statistics
  - ❖ 80-98% of newly allocated are extremely short lived (few million instructions)
  - ❖ 80-98% die before another megabyte has been allocated
  - ❖ Typical programs
- ❖ Tomcat Core (no webapps)
  - ❖ Lots of long lived objects
  - ❖ Still a small memory footprint

# Java Object Heap

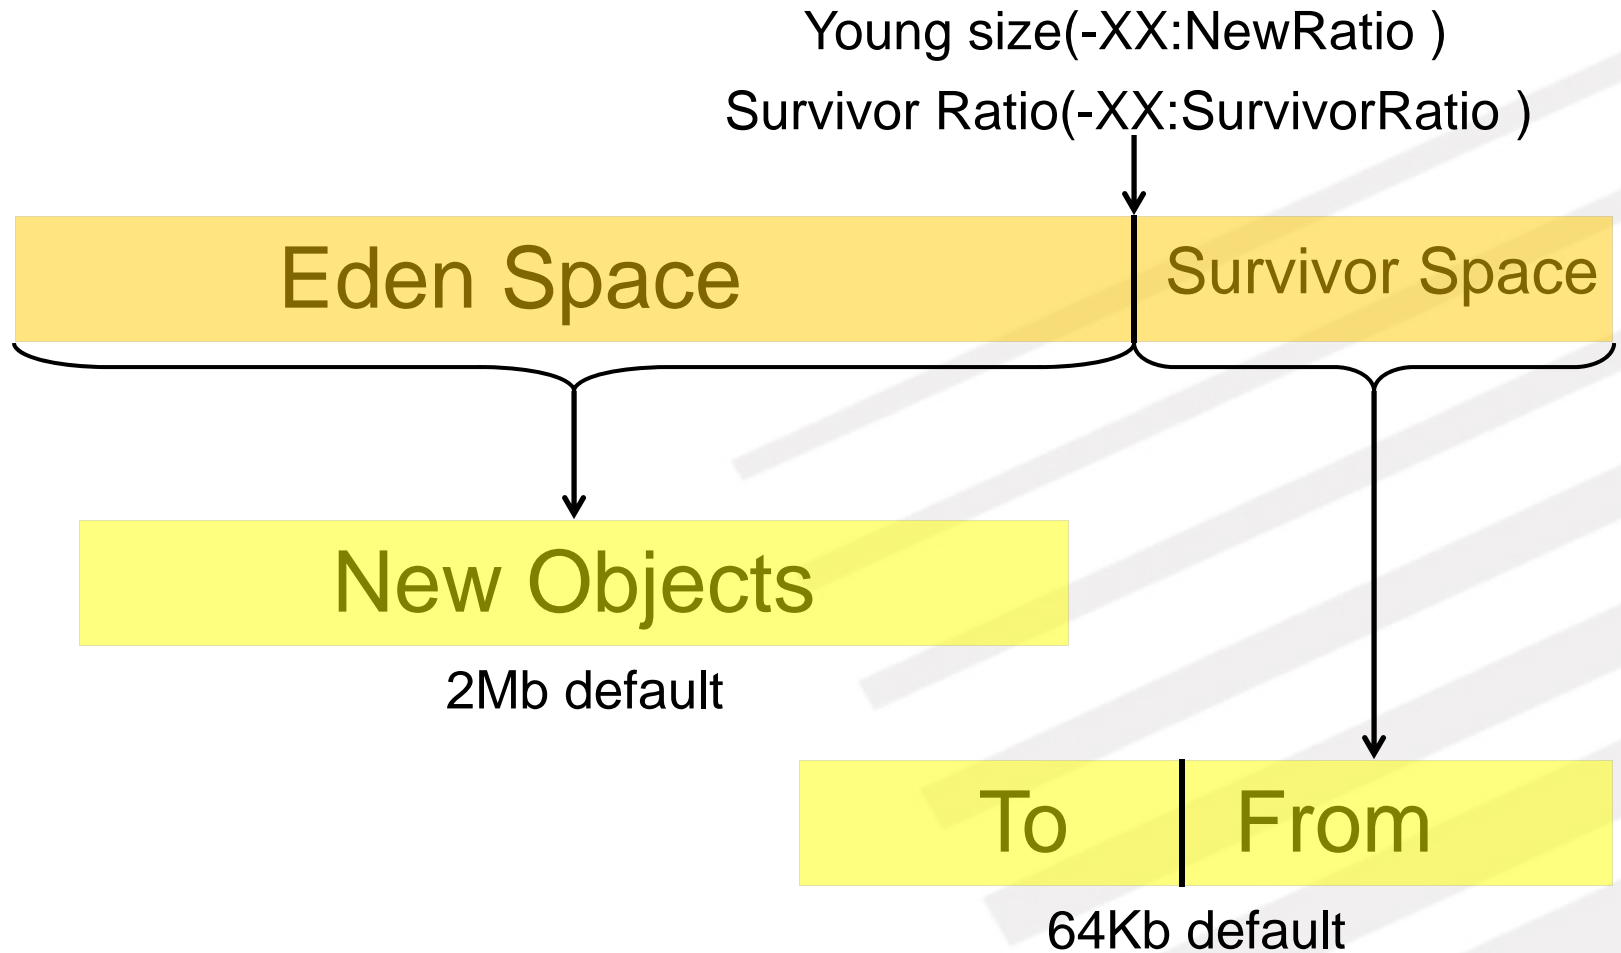


A good size for the YG is 33% of the total heap

# Java Object Heap

- ❖ Young Generation
  - ❖ All new objects are created here
  - ❖ Only moved to Old Gen if they survive one or more **minor GC**
  - ❖ Sized Using
    - ❖ -Xmn – not preferred (fixed value)
    - ❖ -XX:NewRatio=<value> - preferred (dynamic)
- ❖ Survivor Spaces
  - ❖ 2, used during the GC algorithm (minor collections)

# Young Generation





## Gotcha #2

### ❖ Problem

- ❖ Multithreaded apps create new objects at the same time
- ❖ New objects are always created in the EDEN space
- ❖ During object creation, memory is locked
- ❖ On a multi CPU machine (threads run concurrently) there can be contention

# Gotcha #2

- ❖ Solution
  - ❖ Allow each thread to have a private piece of the EDEN space
- ❖ Thread Local Allocation Buffer
  - ❖ -XX:+UseTLAB
  - ❖ -XX:TLABSize=<size in kb>
  - ❖ -XX:+ResizeTLAB
  - ❖ (On by default on multi CPU machines and newer JDK)
- ❖ Analyse TLAB usage
  - ❖ -XX:+PrintTLAB
- ❖ JDK 1.5 and higher (GC ergonomics)
  - ❖ Dynamic sizing algorithm, tuned to each thread

## Tenured Space

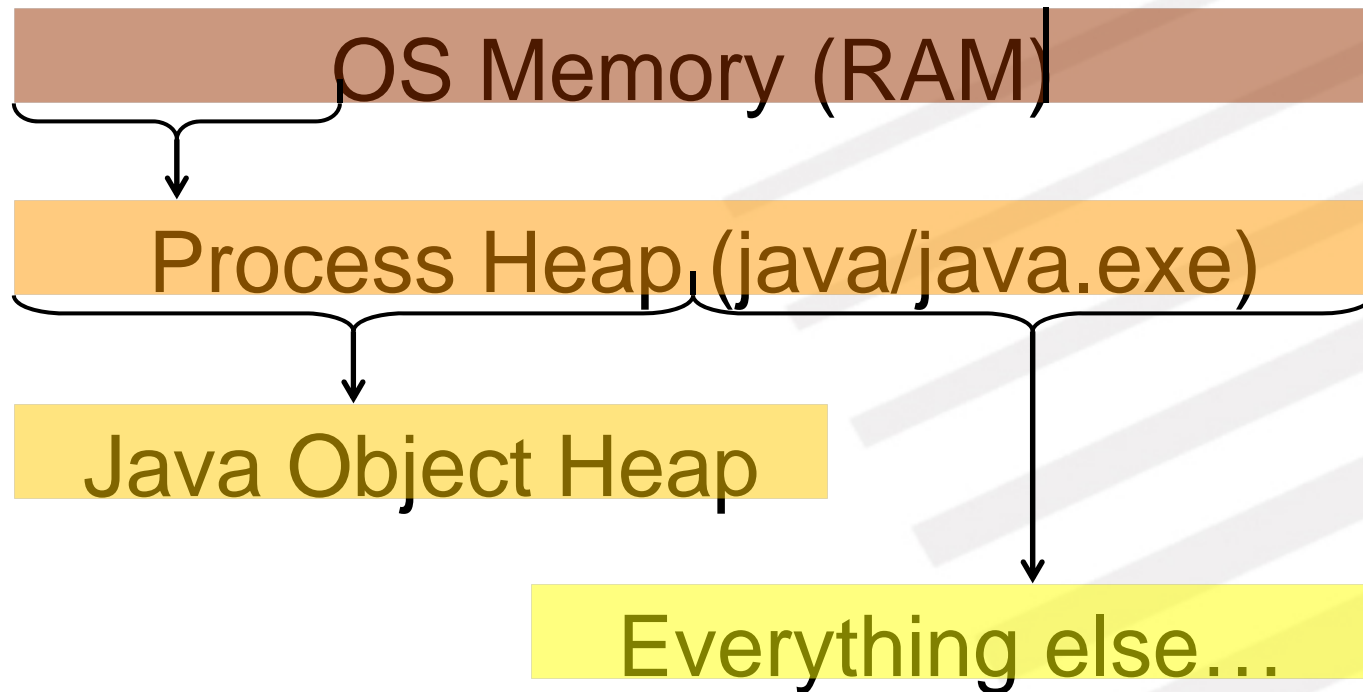
5Mb min 44Mb max (default)

Garbage collection presentation will explain in detail how these spaces are used during the GC process.

# JVM Process Heap

## Java Object Heap

- A handful, but a small part of the story



# JVM Process Heap

## Everything else...

- Permanent Space
- Code Generation
- Socket Buffers
- Thread Stacks
- Direct Memory Space
- JNI Code
- Garbage Collection
- JNI Allocated Memory

# Permanent Space

- ✔ Permanent Generation
  - ✔ Permanent Space (name for it)
  - ✔ 4Mb initial, 64Mb max
  - ✔ Stores classes, methods and other meta data
    - ✔ -XX:PermSize=<value> (initial)
    - ✔ -XX:MaxPermSize=<value> (max)
- ✔ Common OOM for webapp reloads
- ✔ Separate space for pre-historic reasons
  - ✔ Early days of Java, class GC was not common, reduces size of the Java Heap

## Gotcha #3

- ❖ Permanent Space Memory Errors
  - ❖ Too many classes loaded
  - ❖ Classes are not being GC:ed
  - ❖ Unaffected by -Xmx flag
- ❖ Identified by
  - ❖ `java.lang.OutOfMemoryError:`  
PermGen space
- ❖ Many situations, increasing max perm size will help
  - ❖ i.e., no leak, but just not enough memory
  - ❖ Others will require to fix the leak

# JVM Process Heap

## Everything else...

- Permanent Space
- Code Generation
- Socket Buffers
- Thread Stacks
- Direct Memory Space
- JNI Code
- Garbage Collection
- JNI Allocated Memory



# Code Generation

- Converting byte code into native code
- Very rare to cause memory problems
- JVM will most likely crash if it doesn't have enough mem for this operation
  - Never seen it though

# JVM Process Heap

## Everything else...

- Permanent Space
- Code Generation
- Socket Buffers
- Thread Stacks
- Direct Memory Space
- JNI Code
- Garbage Collection
- JNI Allocated Memory

# TCP connections

- ✓ Each connection contains two buffers
  - ✓ Receive buffer ~37k
  - ✓ Send buffer ~25k
- ✓ Configured in Java code
  - ✓ So might not be exposed through applications configuration
- ✓ Usually hit other limits than memory before an error happen
  - ✓ IOException: Too many open files (for example)

# JVM Process Heap

## Everything else...

- Permanent Space
- Code Generation
- Socket Buffers
- Thread Stacks
- Direct Memory Space
- JNI Code
- Garbage Collection
- JNI Allocated Memory

# Thread Stacks

- Each thread has a separate memory space called “thread stack”
- Configured by `-Xss`
- Default value depends on OS/JVM
- As number of threads increase, memory usage increases

# Gotcha #4

- ❖ `java.lang.OutOfMemoryError:`  
unable to create new native thread
- ❖ Solution
  - ❖ Decrease `-Xmx` and/or
  - ❖ Decrease `-Xss`
  - ❖ Or, you have a thread leak, fix the program
- ❖ Gotcha
  - ❖ Increasing `-Xmx` (32bit systems) will leave less room for threads if it is being used, hence the opposite of the solution
  - ❖ Too low `-Xss` value can cause `java.lang.StackOverflowError`

# JVM Process Heap

## Everything else...

- Permanent Space
- Code Generation
- Socket Buffers
- Thread stacks
- Direct Memory Space
- JNI Code
- Garbage Collection
- JNI Allocated Memory

# Direct Memory Space

- Ability to let Java developers map memory outside the Java Object Heap
- `java.nio.ByteBuffer.allocateDirect`
- `java.lang.OutOfMemoryError:`  
Direct buffer memory
- Adjusted by
  - `-XX:MaxDirectMemorySize=<value>`



# JVM Process Heap

## Everything else...

- ✔ Permanent Space
- ✔ Code Generation
- ✔ Socket Buffers
- ✔ Thread stacks
- ✔ Direct Memory Space
- ✔ JNI Code
- ✔ Garbage Collection
- ✔ JNI Allocated Memory

- ✔ Code needs memory
  - ✔ Usually very little
- ✔ JNI programs also allocate memory
  - ✔ Error allocating memory.[NativeMemory.c] (my code)
  - ✔ JVM goes berserk or crashes or if the JNI code can handle it gracefully, you're lucky
- ✔ Linux way of dealing with mem leak
  - ✔ Kill the process!

# JVM Process Heap

## Everything else...

- ✔ Permanent Space
- ✔ Code Generation
- ✔ Socket Buffers
- ✔ Thread stacks
- ✔ Direct Memory Space
- ✔ JNI Code
- ✔ **Garbage Collection**
- ✔ JNI allocated memory

# Garbage Collection

- ✔ Also uses memory
  - ✔ Threads
  - ✔ Memory to store GC info
- ✔ If there isn't enough memory for GC, then the system will not be functioning at all

# GC History

- ✦ First time around 1959 – LISP language
- ✦ The idea
  - ✦ automatic memory cleanup
  - ✦ Easier to write code
  - ✦ Easier to debug
- ✦ What it does
  - ✦ Maps memory in memory
  - ✦ The Java Object Heap is such a map

# Phases of GC

- ✔ Lock it down
  - ✔ All objects that are to take part in the GC must be locked, so that they don't mutate
- ✔ Mark
  - ✔ Iterate through all objects
  - ✔ Mark the “unreachable” as garbage
- ✔ Sweep
  - ✔ Remove all previously marked objects
  - ✔ Reclaim memory

# Early Version of Java


- ❖ Garbage Collector wasn't very well tuned
- ❖ Only one algorithm was available
- ❖ Mark and Sweep entire heap
  - ❖ Takes a very long time
  - ❖ Time spent is dependent on the size of the heap
  - ❖ That is why the "Permanent Space" was invented
    - ❖ And cause un/reloading of classes wasn't very common either
- ❖ Also known as "stop-the-world" gc
  - ❖ The entire JVM is locked down

# Strategies

- ❖ Stop The World
- ❖ Incremental
  - ❖ Time GC with new object creation
  - ❖ If GC runs, suspend new allocation
- ❖ Concurrent/Parallel
  - ❖ Allocation happens at the same time as GC
  - ❖ Very complex locking regimes
  - ❖ Generations/Spaces make it easier
- ❖ CMS stands for
  - ❖ Concurrent
  - ❖ Mark
  - ❖ Sweep



# How It Works



1. New objects are created in the 1<sup>st</sup> collection
2. Copy of Eden objects into 2<sup>nd</sup>
3. Copy of 2<sup>nd</sup> objects remain in Eden or 1<sup>st</sup>
4. Copy from 1<sup>st</sup> to 2<sup>nd</sup>
5. Copy from 2<sup>nd</sup> to the tenured

# How it Works

- One survivor space is always empty
  - Serves as destination for minor collections
- Objects get copied to the tenured space when the 2<sup>nd</sup> survivor space fills up
- Major collections occur when the tenured space fills up
  - Major collections free up Eden and both survivor spaces

# New and Fancy

- Concurrent/Parallel Garbage Collection
- -XX:+UseParNewGC
  - Parallel GC in the New(Young) Generation
- -XX:+UseConcMarkSweepGC
  - Concurrent in the Old generation
- Use these two combined
  - Multi CPU box can take advantage of this

# Sun Recommended

- ▼ GC Settings
  - ▼ -XX:+UseConcMarkSweepGC
  - ▼ -XX:+CMSIncrementalMode
  - ▼ -XX:+CMSIncrementalPacing
  - ▼ -XX:CMSIncrementalDutyCycleMin=0
  - ▼ -XX:+CMSIncrementalDutyCycle=10
  - ▼ -XX:+UseParNewGC
  - ▼ -XX:+CMSPermGenSweepingEnabled
- ▼ To analyze what is going on
  - ▼ -XX:+PrintGCDetails
  - ▼ -XX:+PrintGCTimeStamps
  - ▼ -XX:-TraceClassUnloading

# Minor Notes

- ✔ -XX:+UseParallelGC <> -XX:+UseParNewGC
- ✔ -XX:ParallelGCThreads=<nr of cpu>
  - ✔ Use with ParallelGC setting
- ✔ If you have 4 cpus and 1 JVM
  - ✔ Set value to 4
- ✔ If you have 4 cpus and 2 JVM
  - ✔ Set value to 2
- ✔ If you have 4 cpus and 6 JVM
  - ✔ Set value to 2

# GC Ergonomics

- ✔ Started with JDK 1.5
- ✔ JVM is self trained and GC strategy adapts
- ✔ Rules/Guidelines can be set using command line options
  - ✔ Max pause time goal
    - ✔ The longest pause time to suspend application
  - ✔ Throughput goal
    - ✔ Time spent GC vs. time spent outside GC
- ✔ Not guaranteed

# Out Of Memory Errors

- There is a seamless way to get info
  - `-XX:+HeapDumpOnOutOfMemoryError`
- No performance impact during runtime
- Dumping a `-Xmx512m` heap
  - Create a 512MB .hprof file
  - JVM is “dead” during dumping
  - Restarting JVM during this dump will cause unusable .hprof file

# Gotcha's

- Major collections don't run until tenured is full
- What does that mean?
  - -Xmx1024m
  - Current heap could be 750MB
  - 500MB of "dead" objects
  - If VM is idle, could stay like that for a very long time
  - Wasting 500MB of RAM for an idle JVM



# Monitoring Agents

- ✔ Monitor memory usage
- ✔ If system is idle, force a GC
- ✔ Can be done automatically and with remote agents
- ✔ Example would be:
- ✔ [www.yourkit.com](http://www.yourkit.com)
- ✔ And script the client to take telemetry readings from an embedded agent

# Thank You

- ✔ [fhanik@covalent.net](mailto:fhanik@covalent.net)
- ✔ For support information contact Covalent
  - ✔ [support@covalent.com](mailto:support@covalent.com),  
[sales@covalent.com](mailto:sales@covalent.com)
  - ✔ [www.covalent.com](http://www.covalent.com)
  - ✔ 800/444-1935 or 925/974-8800
- ✔ Question and Answer time!!



***The Enterprise Open Source Support Company***